

Lecture 5: Binary heaps

Sorting algorithms: Heapsort and Quicksort

Binary heaps

What is a binary heap?

- array A of objects with 2 special attributes: $A.length$ and $A.heap_size$.
- it represents a complete binary tree with $A.heap_size$ nodes
 - The tree is completely filled on all levels except possibly the lowest, which is filled from left to right
 - $A.length$ represents the maximum number of nodes of the tree. Therefore, $A.heap_size \leq A.length$
- The index of the **parent**, **left child**, and **right child** of a node with index i are computed as follows:

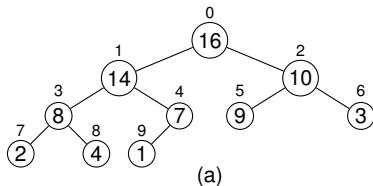
$$parent(i) := \begin{cases} \lfloor (i-1)/2 \rfloor & \text{if } i \neq 0 \\ -1 & \text{if } i = 0 \end{cases}$$

$$left(i) := 2 \cdot i + 1$$

$$right(i) := 2 \cdot i + 2$$

- The **heap property** must hold: $A[parent(i)] \geq A[i]$ for all $i \neq 0$.

Binary heaps: Example



0	1	2	3	4	5	6	7	8	9
16	14	10	8	7	9	3	2	4	1

(b)

A heap viewed as **(a) a binary tree** and **(b) an array**. The number within the circle at each node in the tree is the value stored at that node. The number next to a node is the corresponding index in the array.

AUXILIARY NOTIONS

- **height of a node** in a tree := number of edges from the tree to a leaf.
- **height of the tree** := height of the root of the tree.

- The height of a binary heap is $\Theta(\log_2(n))$ – obvious.
- FIND / INSERT / REMOVE operations in binary heaps take $O(\log_2(n))$ time – we shall prove this.
- We are interested in the efficient implementation of:
 - 1 HEAPIFY(A, i)
 - 2 BUILDHEAP(A)
 - 3 HEAPSORT(A)
 - 4 EXTRACTMAX(A)
 - 5 INSERT(A, key)

The purpose of these procedures will be explained later.

HEAPIFY(A, i)

- Takes as input an array A and an index i , such that
 - the subtrees rooted at $left(i)$ and $right(i)$ are binary heaps.
 - The subtree rooted at i may not be a binary heap, because $A[i]$ is smaller than its children.
- Rearranges the elements of A by letting $A[i]$ "float down" so that the subtree rooted at index i becomes a binary heap.

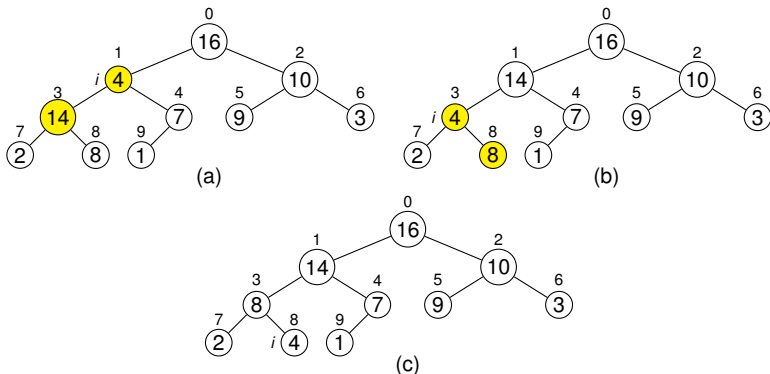
Thus, the purpose of HEAPIFY is to maintain the heap property of an array of values.

HEAPIFY(A, i)

```
HEAPIFY( $A, i$ )  
  1  $l := \text{left}(i)$   
  2  $r := \text{right}(i)$   
  3 if  $l < A.\text{heap\_size}$  and  $A[l] > A[i]$   
  4    $\text{largest} := l$   
  5 else  $\text{largest} := i$   
  6 if  $r < A.\text{heap\_size}$  and  $A[r] > A[\text{largest}]$   
  7    $\text{largest} := r$   
  8 if  $\text{largest} \neq i$   
  9   exchange  $A[i] \leftrightarrow A[\text{largest}]$   
 10  HEAPIFY( $A, \text{largest}$ )
```

Example

The action of $\text{HEAPIFY}(A, 1)$, where $A.\text{heap_size} = 10$. Configuration **(a)** lacks heap property at index 1. The heap property for index 1 is restored in **(b)** by exchanging $A[1]$ with $A[3]$, which destroys the heap property for index 3. The recursive call $\text{HEAPIFY}(A, 3)$ sets $i = 3$, swaps $A[3] \leftrightarrow A[8]$ as shown in **(c)**, and the recursive call $\text{HEAPIFY}(A, 8)$ yields no further change to the data structure.



Properties of HEAPIFY

- The running time complexity of $\text{HEAPIFY}(A, i)$ is $O(h)$, where h is the height of node with index i .
- \Rightarrow In general, the running time of $\text{HEAPIFY}(A, i)$ is $O(\log_2(n))$.
- For a proof, check the references.

Building a binary heap

BUILDHEAP(A)

- Rearranges the elements of an array A , to have the binary heap property.
- The rearrangement is achieved by successive runs of HEAPIFY(A, i)

BUILDHEAP(A)

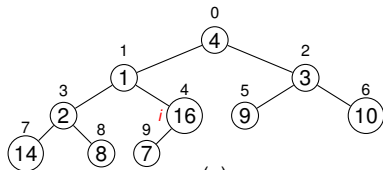
```
1   $heap\_size(A) := A.length$   
2  for  $i := \lfloor (A.length - 1)/2 \rfloor$  downto 0  
3      HEAPIFY( $A, i$ )
```

Remarks

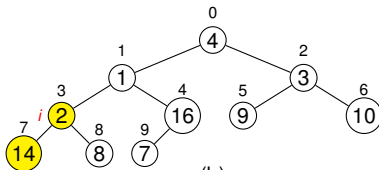
- The order in which the nodes are processed guarantees that the subtrees rooted at children of a node i are heaps before HEAPIFY is run at that node.
- There are $O(n)$ calls of HEAPIFY(A, i), which has time complexity $O(\log_2 n) \Rightarrow$ time complexity $O(n \log_2 n)$.
- Tighter bound of the total runtime of step 3: $O(n)$ (see refs.)

Example

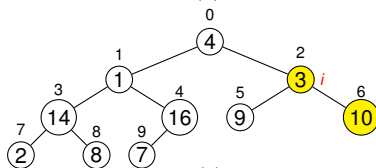
BUILDHEAP(A) for $A=\{4,1,3,2,16,9,10,14,8,7\}$.



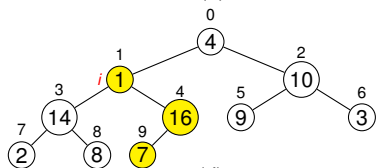
(a)



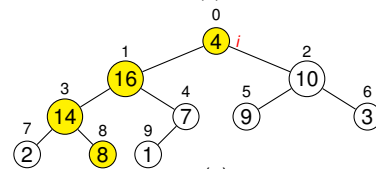
(b)



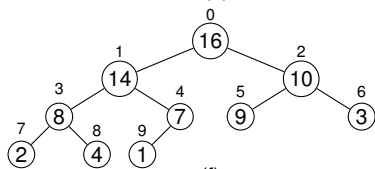
(c)



(d)



(e)



(f)

The Heapsort algorithm

HEAPSORT(A) rearranges the elements of an array A in ascending order, using the following method:

- 1 Call **BUILDHEAP(A)** \Rightarrow a heap on the elements of the array $A[0..n-1]$
- 2 $A[0]$ is the maximum element of A
 - ▷ exchange $A[0] \leftrightarrow A[n-1]$, to place $A[0]$ into its correct final position.
- 3 Discard $A[n-1]$ from the heap by decrementing $A.heap_size$. We still have to sort $A[0..n-2]$
 - $A[0..n-2]$ is *almost* a binary heap: 0 is the only index that may violate the heap property.
 - We run **HEAPIFY($A, 0$)** to rearrange $A[0..n-2]$ into binary heap.
 - The Heapsort algorithm **repeats this process** for the heap of size $n-1$ down to a heap of size 2.

Heapsort

HEAPSORT(A)

1 BUILDHEAP(A)

2 **for** $i := A.length - 1$ **downto** 1

3 exchange $A[0] \leftrightarrow A[i]$

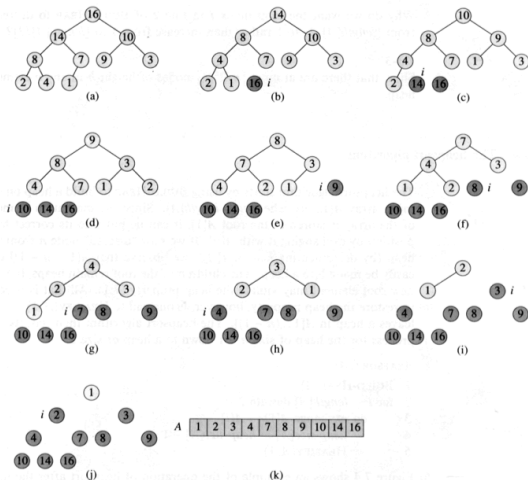
4 $A.heap_size := A.heap_size - 1$

5 HEAPIFY($A, 0$)

TIME COMPLEXITY ANALYSIS

- BUILDHEAP(A) takes $O(n)$ time.
 - There are $n - 1$ calls to HEAPIFY($A, 0$), and each one takes $O(\log_2 n)$ time.
- ⇒ HEAPSORT(A) takes $O(n \log_2 n)$ time, where $n = A.length$.

Heapsort – running example



(a) The heap data structure just after it has been built by BUILDHEAP. **(b)–(j)** The heap just after each call of HEAPIFY in line 5. The value of i at that time is shown. Only lightly shaded nodes remain in the heap. **(k)** The resulting sorted array A .

Priority queues

A **priority queue** is a data structure for maintaining a set S of elements, each with an associated value called a **key**. It is intended to support efficient execution of the following operations:

- **INSERT(S, x)**: inserts the element x into a set S . We denote this operation by $S := S \cup \{x\}$.
- **MAXIMUM(S)**: returns the element of S with the largest key.
- **EXTRACTMAX(S)**: removes and returns the element of S with the largest key.

Applications of priority queues

- Job scheduling on a shared resource
 - The queue keeps track of jobs to be performed, and their relative priorities.
 - When a job is finished or interrupted, the highest-priority job is selected from the queue, using **EXTRACTMAX**
 - New jobs can be added at any time using **INSERT**
- Event-driven simulation: time of event occurrence serves as its key.

Priority queues

Can be implemented efficiently using binary heaps.

EXTRACTMAX(A)

```
1  if  $A.heap\_size < 1$ 
2    error "heap underflow"
3   $max := A[0]$ 
4   $A[0] := A[A.heap\_size - 1]$ 
5   $A.heap\_size := A.heap\_size - 1$ 
6  HEAPIFY( $A, 0$ )
7  return  $max$ 
```

Running time analysis

- HEAPIFY($A, 0$) takes $O(\log_2 n)$ time
⇒ EXTRACTMAX(A) takes $O(\log_2 n)$ time.

Priority queues

INSERT(A, key)

INSERT(A, key) inserts a node into a binary heap A :

- First, it expands the heap by adding a new leaf to the tree.
- Then, it traverses a path from this leaf toward the root, to find a proper place for the new element.

INSERT(A, key)

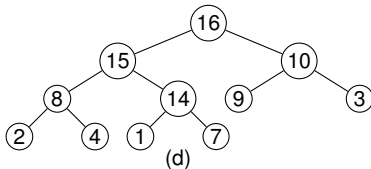
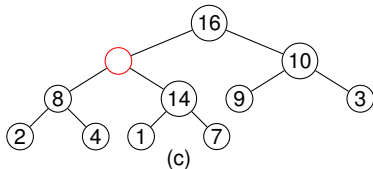
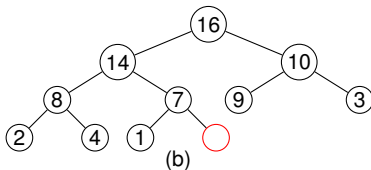
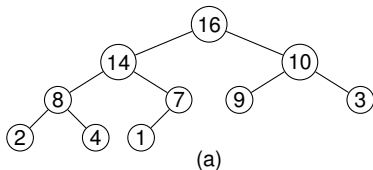
```
1   $A.heap\_size := A.heap\_size + 1$ 
2   $i := A.heap\_size - 1$ 
3  while  $i > 0$  and  $A[parent(i)] < key$ 
4       $A[i] := A[parent(i)]$ 
5       $i := parent(i)$ 
6   $A[i] := key$ 
```

Running time analysis

- The path traced from the new leaf to the root has length $O(\log_2 n) \Rightarrow$ HEAPINSERT(A, key) takes $O(\log_2 n)$ time, where $n = A.heap_size$.

Priority queues

INSERT(A, key) illustrated



- (a)** The heap before we insert a node with key 15. **(b)** A new leaf is added to the tree. **(c)** Values on the path from the new leaf to the root are copied down until a place for the key 15 is found. **(d)** Key 15 is inserted into the tree.

Quicksort

Properties

- Sorting algorithm with worst-case running time $\Theta(n^2)$ on an input array of n numbers.
- Very efficient on average: $\Theta(n \log n)$
- Often, the best practical choice for sorting

Quicksort

Description of the algorithm

3-step divide-and-conquer algorithm for sorting a subarray $A[p..r]$

Divide: The subarray $A[p..r]$ is partitioned (rearranged) into two nonempty subarrays $A[p..q]$, $A[q + 1..r]$ such that

- The elements of $A[p..q]$ are smaller than the elements of $A[q + 1..r]$

The index q is computed as part of this partitioning procedure.

Conquer: The subarrays $A[p..q]$ and $A[q + 1..r]$ are sorted by recursive calls to quicksort.

Combine: Since the subarrays are sorted in place, no work is needed to combine them: the entire array $A[p..r]$ is now sorted.

Quicksort

Pseudocode

QUICKSORT(A, p, r)

1. **if** $p < r$
2. $q \leftarrow \text{PARTITION}(A, p, r)$
3. QUICKSORT(A, p, q)
4. QUICKSORT($A, q + 1, r$)

Partitioning the array

PARTITION(A, p, r)

```
1   $x \leftarrow A[p]$ 
2   $i \leftarrow p - 1$ 
3   $j \leftarrow r + 1$ 
4  while TRUE
5      do repeat  $j \leftarrow j - 1$ 
6          until  $A[j] \leq x$ 
7      repeat  $i \leftarrow i + 1$ 
8          until  $A[i] \geq x$ 
9      if  $i < j$ 
10         then exchange  $A[i] \leftrightarrow A[j]$ 
11         else return  $j$ 
```

Quicksort

How does PARTITION work?

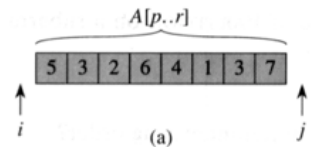
- ▶ Element $x = A[p]$ from $A[p..r]$ is selected as **pivot** around which to partition $A[p..r]$.
- ▶ The **while** loop grows two regions $A[p..i]$ and $A[j..r]$ from the top and bottom of $A[p..r]$, respectively, such that
 - Every element in $A[p..i]$ is less than or equal to x .
 - Every element in $A[j..r]$ is greater than or equal to x .

Initially, $i = p - 1$ and $j = r + 1$, so the two regions are empty.

- ▶ Within the **while** loop, index j is decremented and index i is incremented, in lines 5-8, until $A[i] \geq x \geq A[j]$.
 - By exchanging $A[i]$ and $A[j]$, the two regions can be extended.
- ▶ The **while** loop repeats until $i \geq j$, at which point the entire array $A[p..r]$ has been partitioned into two subarrays $A[p..q]$ and $A[q + 1..r]$ where $p \leq q < r$, such that all elements in $A[p..q]$ are smaller than or equal to any element in $A[q + 1..r]$.
- ▶ The value $q = j$ is returned at the end of the procedure.

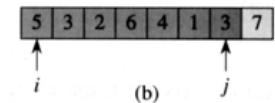
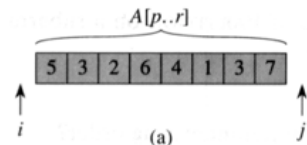
Quicksort

Example of how PARTITION works



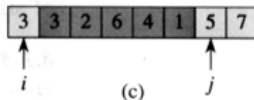
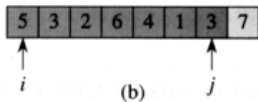
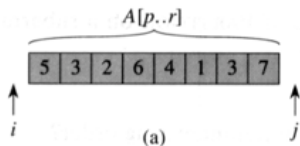
Quicksort

Example of how PARTITION works



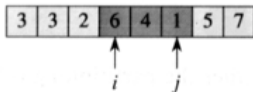
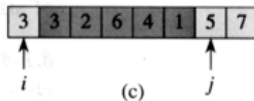
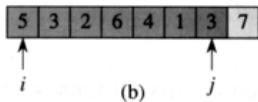
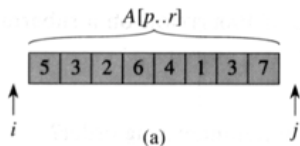
Quicksort

Example of how PARTITION works



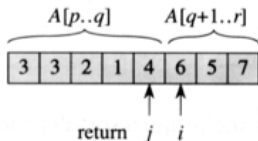
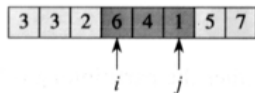
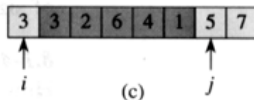
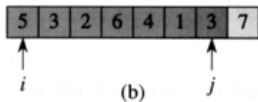
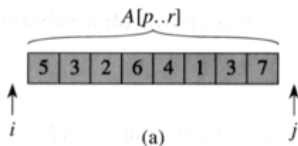
Quicksort

Example of how PARTITION works



Quicksort

Example of how PARTITION works



Quicksort

Complexity analysis

- The running time of PARTITION on an array $A[p..r]$ is $\Theta(r - p + 1)$.
- Worst case behavior happens when the partitioning always produces one partition with 1 element, and the other with all the rest. In this case:
 - Partitioning an array of size n takes $\Theta(n)$ time and $T(1) = \Theta(1)$.
 - The recurrence relation is $T(n) = T(n - 1) + \Theta(n - 1) = \dots = \sum_{k=1}^n \Theta(k) = \Theta(\sum_{k=1}^n k) = \Theta(n^2)$. \Rightarrow in the worst case, the running time is $\Theta(n^2)$.
- Best case is when the partitioning produces regions of equal size \Rightarrow the recurrence relation $T(n) = 2 T(n/2) + \Theta(n)$.
 - $\Rightarrow T(n) = \Theta(n \log n)$
(Cf. the Master Theorem)

Chapters 7 (Heapsort) and 8 (Quicksort) from the book

- Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest. *Introduction to Algorithms*. McGraw Hill, 2000.